



CMPT 713: Natural Language Processing

Word Embeddings - Word2Vec

Spring 2023
2023-01-26

Adapted from slides from Dan Jurafsky, Chris Manning, Danqi Chen and Karthik Narasimhan

Representing words by their context

Distributional hypothesis: words that occur in similar contexts tend to have similar meanings



J.R.Firth 1957

- “You shall know a word by the company it keeps”
- One of the most successful ideas of modern statistical NLP!

*...government debt problems turning into **banking** crises as happened in 2009...*

*...saying that Europe needs unified **banking** regulation to replace the hodgepodge...*

*...India has just given its **banking** system a shot in the arm...*

These **context words** will represent **banking**.

Word Vectors

- One-hot vectors

[illegible]
$$\text{motel} = [0 \ 0 \ 0 \ \mathbf{1} \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$$

- Represent words by their context

word-word (term-context) co-occurrence matrix

context

(other words in the span
around the target word)

	aardvark	computer	data	pinch	result	sugar	...
apricot	0	0	0	1	0	1	
pineapple	0	0	0	1	0	1	
digital	0	2	1	0	1	0	
information	0	1	6	0	4	0	

$|V| \times |V|$ matrix

sugar, a sliced lemon, a tablespoonful of **apricot** jam, a pinch each of, their enjoyment. Cautiously she sampled her first **pineapple** and another fruit whose taste she likened well suited to programming on the digital **computer**. In finding the optimal R-stage policy from for the purpose of gathering data and **information** necessary for the study authorized in the

Sparse vs dense vectors

Vectors we get from word-word (term-context) co-occurrence matrix are

- **long** (length $|V| = 20,000$ to $50,000$)
- **sparse** (most elements are zero)

True for both one-hot, tf-idf and PPMI vectors

Alternative: we want to represent words as

- **short** (50-300 dimensional)
- **dense** (real-valued) vectors
- The focus of this lecture
- The basis of all the modern NLP systems

Dense vectors

$$\text{employees} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 10.109 \\ -0.542 \\ 0.349 \\ 0.271 \\ 0.487 \end{pmatrix}$$


short + dense



Why dense vectors?

- Short vectors are easier to use as features in ML systems
- Dense vectors may generalize better than storing explicit counts
- They do better at capturing synonymy
 - w_1 co-occurs with “car”, w_2 co-occurs with “automobile”

- Different methods for getting dense vectors:
 - Singular value decomposition (SVD)
 - word2vec and friends: “learn” the vectors!



Count based method
(known since the 1990s)

What is wrong with using SVD?

- Computational complexity is high: $O(|V|^3)$
- Cannot be trained as part of a larger model
 - Not a component that can be part of a larger neural network
- Cannot be trained discriminatively for a particular task

Why dense vectors?

- Short vectors are easier to use as features in ML systems
- Dense vectors may generalize better than storing explicit counts
- They do better at capturing synonymy
 - w_1 co-occurs with “car”, w_2 co-occurs with “automobile”
- Different methods for getting dense vectors:
 - Singular value decomposition (SVD)
 - word2vec and friends: “learn” the vectors! ← Typically what is referred to when people say “word embeddings”
Don’t count, predict! (Baroni et al, 2014)

Word2vec and friends

Download pretrained word embeddings

Google

Word2vec (Mikolov et al.)

<https://code.google.com/archive/p/word2vec/>

facebook

Fasttext <http://www.fasttext.cc/>



Glove (Pennington, Socher, Manning)

<http://nlp.stanford.edu/projects/glove/>



Word2Vec

- Popular embedding method
- Very fast to train
- Idea: **predict** rather than **count**



Original word2vec formulation: Efficient Estimation of Word Representations in Vector Space (Mikolov et al, 2013a)

Negative sampling: Distributed Representations of Words and Phrases and their Compositionality (Mikolov et al, 2013b)


Word2Vec

- Instead of counting how often each word w occurs near “apricot”
- Train a classifier on a **prediction** task:
 - Is w likely to show up near “apricot”?
 - There are different ways to formulate this prediction task
 - We will look at “Skip-gram with negative sampling”
- We don’t actually care about this task (pretext task)
 - But we’ll take the learned classifier **weights** as **the word embeddings**

Word2Vec

Insight: use running text as implicitly supervised training data!

- A word s near apricot
 - Act as gold “correct answer” to the question “Is word w likely to show up near apricot?”
- No need for hand-labeled supervision
- The idea comes from neural language modeling
 - Bengio et al (2003)
 - Collobert et al (2011)



(Bengio et al, 2003): A Neural Probabilistic Language Model

Word2vec


- Input: a large text corpora V, d
 - V : a pre-defined vocabulary
 - d : dimension of word vectors (e.g. 300)
 - Text corpora (words $w_1, \dots, w_T$)
 - Wikipedia + Gigaword 5: 6B
 - Twitter: 27B
 - Common Crawl: 840B
- Output: $f : V \rightarrow \mathbb{R}^d$

$$v_{\text{cat}} = \begin{pmatrix} -0.224 \\ 0.130 \\ -0.290 \\ 0.276 \end{pmatrix} \quad v_{\text{dog}} = \begin{pmatrix} -0.124 \\ 0.430 \\ -0.200 \\ 0.329 \end{pmatrix}$$
$$v_{\text{the}} = \begin{pmatrix} 0.234 \\ 0.266 \\ 0.239 \\ -0.199 \end{pmatrix} \quad v_{\text{language}} = \begin{pmatrix} 0.290 \\ -0.441 \\ 0.762 \\ 0.982 \end{pmatrix}$$


Learn f by training classifiers to **predict** words and take learned weights as word vectors.

Word2vec

Predict **center** word from context words Predict **context** words from center word




Continuous Bag of Words (CBOW)



Skip-grams

Skip-gram

- The idea: we want to use words to **predict** their context words
- Context: a fixed window of size $2m$ ($m = 2$ in this example)



Skip-gram: Basic Setup

Let's represent words as vectors of some length (say 300), randomly initialized.

So we start with $300 * V$ random parameters

Over the entire training set, we'd like to adjust those word vectors such that we

- Predict the probability distribution for how likely a word c is to be a context word of a target word t

$$P(c | t)$$

Skip-gram

Training sentence:

... lemon, a tablespoon of apricot jam a pinch ...

c1 c2 t c3 c4

Training data

Training data: input/output pairs centering on apricot
assume a +/- 2 word window

Given a tuple (t, c) = target, context

(apricot, jam)

(apricot, aardvark)

Goal

Return probability that c is a real context word:

$$P(c | t)$$

Probability distribution
over vocabulary

Skip-gram: objective function

- For each position $t = 1, 2, \dots, T$, predict context words within context size m , given center word w_j :

$$\mathcal{L}(\theta) = \prod_{t=1}^T \prod_{-m \leq j \leq m, j \neq 0} P(w_{t+j} \mid w_t; \theta)$$

all the parameters to be optimized

Subscript refer to position in corpus
 w_{t+j} refer to word in vocabulary V

- The objective function $J(\theta)$ is the (average) negative log likelihood:

$$J(\theta) = -\frac{1}{T} \log \mathcal{L}(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} \mid w_t; \theta)$$

How to define $P(w_{t+j} \mid w_t; \theta)$?

- We have two sets of vectors for each word in the vocabulary

$\mathbf{u}_i \in \mathbb{R}^d$: embedding for target word i

$\mathbf{v}_{i'} \in \mathbb{R}^d$: embedding for context word i'

- Use inner product $\mathbf{u}_i \cdot \mathbf{v}_{i'}$ to measure how likely word i appears with context word i' , the larger the better

$$P(w_{t+j} \mid w_t) = \frac{\exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_{w_{t+j}})}{\sum_{k \in V} \exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_k)}$$

“softmax” we learned before!
← Normalized over entire vocabulary

$\theta = \{\{\mathbf{u}_k\}, \{\mathbf{v}_k\}\}$ are all the parameters in this model!

Q: Why two sets of vectors?

How to train the model

Objective function: average log likelihood

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m, j \neq 0} \log P(w_{t+j} \mid w_t; \theta)$$

Parameters:

Q: How many parameters are in total?

$$\theta = \{\{\mathbf{u}_k\}, \{\mathbf{v}_k\}\}$$

Note: because we need to learn both set of parameters, this is a non-convex objective function

We can apply stochastic gradient descent (SGD)!

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla_{\theta} J(\theta)$$

Need to compute: $\nabla_{\theta} J(\theta) = ?$

Computing the gradients

Consider one pair of target/context words (t, c):

$$y = -\log \left(\frac{\exp(\mathbf{u}_t \cdot \mathbf{v}_c)}{\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k)} \right) = -\mathbf{u}_t \cdot \mathbf{v}_c + \log \left(\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k) \right) \quad \forall k \in V$$

$$\begin{aligned} \frac{\partial y}{\partial \mathbf{u}_t} &= \frac{\partial(-\mathbf{u}_t \cdot \mathbf{v}_c + \log(\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k)))}{\partial \mathbf{u}_t} \\ &= -\mathbf{v}_c + \frac{\sum_{k \in V} \frac{\partial \exp(\mathbf{u}_t \cdot \mathbf{v}_k)}{\partial \mathbf{u}_t}}{\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k)} \\ &= -\mathbf{v}_c + \frac{\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k) \mathbf{v}_k}{\sum_{k' \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_{k'})} \\ &= -\mathbf{v}_c + \sum_{k \in V} \frac{\exp(\mathbf{u}_t \cdot \mathbf{v}_k)}{\sum_{k' \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_{k'})} \mathbf{v}_k \\ &= -\mathbf{v}_c + \sum_{k \in V} P(k|t) \mathbf{v}_k \end{aligned}$$

$$\begin{aligned} \frac{\partial y}{\partial \mathbf{v}_k} &= \frac{\partial(-\mathbf{u}_t \cdot \mathbf{v}_c + \log(\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k)))}{\partial \mathbf{v}_k} \\ &= -1[k = c] \mathbf{u}_t + P(k|t) \mathbf{u}_t \end{aligned}$$

We can pull $\sum_{k \in V}$ out and push the denominator inside the sum

Note we defined $P(w_{t+j} \mid w_t) = \frac{\exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_{w_{t+j}})}{\sum_{k \in V} \exp(\mathbf{u}_{w_t} \cdot \mathbf{v}_k)}$

Make sure you know how to do this!

Putting it together

- Input: text corpus, context size m , embedding size d , V
- Initialize $\mathbf{u}_i, \mathbf{v}_i$ randomly
- Walk through the training corpus and collect training data (t, c) :

- Update $\mathbf{u}_t \leftarrow \mathbf{u}_t - \eta \frac{\partial y}{\partial \mathbf{u}_t}$
- Update $\mathbf{v}_k \leftarrow \mathbf{v}_k - \eta \frac{\partial y}{\partial \mathbf{v}_k}, \forall k \in V$

Any issues?

Problem with Naive softmax

Problem: every time you get one pair of (t, c) , you need to update $\mathbf{u}_t$ for each $\mathbf{v}_k$, and update $\mathbf{v}_k$ for all the words in the vocabulary! It is very computationally expensive.

$$\frac{\partial y}{\partial \mathbf{u}_t} = -\mathbf{v}_c + \sum_{k \in V} P(k|t) \mathbf{v}_k$$

$$\frac{\partial y}{\partial \mathbf{v}_k} = -1[k = c] \mathbf{u}_t + P(k|t) \mathbf{u}_t \quad \forall k \in V$$

Negative sampling: instead of considering all the words in V , let's randomly sample K (5-20) negative examples.

softmax:
$$y = -\log \left(\frac{\exp(\mathbf{u}_t \cdot \mathbf{v}_c)}{\sum_{k \in V} \exp(\mathbf{u}_t \cdot \mathbf{v}_k)} \right)$$

NS:
$$y = -\log(\sigma(\mathbf{u}_t \cdot \mathbf{v}_c)) - \sum_{i=1}^K E_{j \sim P(w)} \log(\sigma(-\mathbf{u}_t \cdot \mathbf{v}_j))$$

Skip-gram with negative sampling (SGNS)

Key idea: convert the $|V|$ -way classification problem into a binary classification problem

$$y = -\log(\sigma(\mathbf{u}_t \cdot \mathbf{v}_c)) - \sum_{i=1}^K E_{j \sim P(w)} \log(\sigma(-\mathbf{u}_t \cdot \mathbf{v}_j))$$

Every time we get a pair of (t, c) words, predict if they co-occur together or not.
(don't try to predict the correct c from amongst all the words in the vocabulary)


positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

negative examples -

t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



Similar to training a **logistic regression** for binary classification (but need to optimize $\mathbf{u}$ and $\mathbf{v}$ together)

$$P(y = 1 \mid t, c) = \sigma(\mathbf{u}_t \cdot \mathbf{v}_c)$$

Skip-gram with negative sampling (SGNS)

Basic formula:

1. Treat the target word and a neighboring context word as positive examples.
2. Randomly sample other words in the lexicon to get negative samples
3. Train a classifier to distinguish those two cases
4. Use the weights as the embeddings

SGNS: Training the classifier

Iterative process

Let's represent words as vectors of some length (say 300), randomly initialized.

So we start with $300 * V$ random parameters

Over the entire training set, adjust those word vectors

- Maximize the similarity of the **target word**, **context word** pairs (t,c) drawn from the positive data
- Minimize the similarity of the (t,c) pairs drawn from the negative data.

Skip-gram with negative sampling

Training sentence:

... lemon, a tablespoon of apricot jam a pinch ...

c1 c2 t c3 c4

Training data

Training data: input/output pairs centering on apricot
assume a +/- 2 word window

Given a tuple (t, c) = target, context

(apricot, jam)

(apricot, aardvark)

Goal

Return probability that c is a real context word:

$$P(+ | t, c)$$

$$P(- | t, c) = 1 - P(+ | t, c)$$

SGNS Training

Training sentence:

... lemon, a tablespoon of apricot jam a pinch ...
 c1 c2 t c3 c4

positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	preserves
apricot	or

- For each positive example, we'll create K negative examples.
- Using noise words
- Any random word that isn't t

Choosing noise words

Could pick w according to their unigram frequency $P(w)$

More common to chose them according to $P_\alpha(w)$

$$P_\alpha(w) = \frac{\text{count}(w)^\alpha}{\sum_w \text{count}(w)^\alpha}$$

$\alpha = 3/4$ works well because it gives rare noise words slightly higher probability

To show this, imagine two events $p(a) = 0.99$ and $p(b) = 0.01$:

$$P_\alpha(a) = \frac{.99^{.75}}{.99^{.75} + .01^{.75}} = .97$$

$$P_\alpha(b) = \frac{.01^{.75}}{.99^{.75} + .01^{.75}} = .03$$

SGNS: objective function

We want to maximize...

$$\sum_{(t,c) \in +} \log P(+|t, c) + \sum_{(t,c) \in -} \log P(-|t, c)$$

Maximize the + label for the pairs from the **positive** training data,
and the – label for the pairs sample from the **negative** data.

Focusing on one target word t :

$$\begin{aligned} L(\boldsymbol{\theta}) &= \log P(+|t, c) + \sum_{i=1} \log P(-|t, n_i) \\ &= \log \sigma(c \cdot t) + \sum_{i=1}^k \log \sigma(-n_i \cdot t) \\ &= \log \frac{1}{1 + e^{-c \cdot t}} + \sum_{i=1}^k \log \frac{1}{1 + e^{n_i \cdot t}} \end{aligned}$$

Skip-gram with negative sampling (SGNS)

Key idea: convert the $|V|$ -way classification problem into a binary classification problem

$$y = -\log(\sigma(\mathbf{u}_t \cdot \mathbf{v}_c)) - \sum_{i=1}^K E_{j \sim P(w)} \log(\sigma(-\mathbf{u}_t \cdot \mathbf{v}_j))$$

Every time we get a pair of (t, c) words, predict if they co-occur together or not.
(don't try to predict the correct c from amongst all the words in the vocabulary)


positive examples +

t	c
apricot	tablespoon
apricot	of
apricot	jam
apricot	a

negative examples -

t	c	t	c
apricot	aardvark	apricot	seven
apricot	my	apricot	forever
apricot	where	apricot	dear
apricot	coaxial	apricot	if

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$



Similar to training a **logistic regression** for binary classification (but need to optimize $\mathbf{u}$ and $\mathbf{v}$ together)

$$P(y = 1 \mid t, c) = \sigma(\mathbf{u}_t \cdot \mathbf{v}_c)$$

<http://jalammar.github.io/illustrated-word2vec/>

Relationship of PMI and Word2Vec

- Word2Vec Skipgram with negative sampling (SGNS) implicitly factorizes word-context PMI matrix $\mathbf{M}^{\text{PMI}}$

$$\mathbf{W} \cdot \mathbf{C}^T \approx \mathbf{M}^{\text{PMI}} - \log k$$

$$\vec{w} \cdot \vec{c} \approx \text{PMI}(w, c) - \log k$$

$\mathbf{W}$ is the matrix of word embeddings

$\mathbf{C}$ is the matrix of context embeddings


k is the number of negative samples,
and the samples follows the unigram distribution

- Other differences in factorization
 - SGNS factorizes into two unconstrained matrices (vs two orthonormal and one diagonal matrix for SVD)
 - The loss used for the factorizing is different
 - SVD - Frobenius norm (Euclidean norm / entrywise L2)
 - Word2Vec SGNS - weighted logistic loss

[Neural Word Embedding as Implicit Matrix Factorization, Levy and Goldberg, 2014]

[Improving Distributional Similarity with Lessons Learned from Word Embeddings, Levy, Goldberg, and Dagan, 2015]

Continuous Bag of Words (CBOW)



$$L(\theta) = \prod_{t=1}^T P(w_t \mid \{w_{t+j}\}, -m \leq j \leq m, j \neq 0)$$

$$\bar{\mathbf{v}}_t = \frac{1}{2m} \sum_{-m \leq j \leq m, j \neq 0} \mathbf{v}_{t+j}$$

$$P(w_t \mid \{w_{t+j}\}) = \frac{\exp(\mathbf{u}_{w_t} \cdot \bar{\mathbf{v}}_t)}{\sum_{k \in V} \exp(\mathbf{u}_k \cdot \bar{\mathbf{v}}_t)}$$

Skip-gram vs CBOW


- CBOW is comparatively faster to train than skip-grams and better for frequently occurring words
- Skip-gram is slower but works well for smaller amount of data and works well for less frequently occurring words
- CBOW is a simpler problem than Skip-gram because in CBOW we just need to predict the one center word given many context words

GloVe: **G**lobal **V**ectors

- Let's take the global co-occurrence statistics: $X_{i,j}$
- Try to learn word vectors to predict the co-occurrence counts (using L2 loss)
- Function f to weight loss by frequency of words (from 0 to 1)

$$J = \sum_{i,j=1}^{|V|} f(X_{ij})(w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log X_{ij})^2$$

- Final word vector: $w_i + \tilde{w}_j$
- Training faster
- Scalable to very large corpora



(Pennington et al, 2014): GloVe: Global Vectors for Word Representation

GloVe: **G**lobal **V**ectors

Nearest words to
frog:

1. frogs
2. toad
3. litoria
4. leptodactylidae
5. rana
6. lizard
7. eleutherodactylus



litoria



leptodactylidae



rana



eleutherodactylus

(Pennington et al, 2014): GloVe: Global Vectors for Word Representation

FastText: Sub-Word Embeddings

- Similar as Skip-gram, but break words into n-grams with $n = 3$ to 6

where: 3-grams: <wh, whe, her, ere, re>

4-grams: <whe, wher, here, ere>

5-grams: <wher, where, here>

6-grams: <where, where>

Note: All the embeddings that we have learned are also called “static word embeddings”: there is one fixed vector for every word in the vocabulary.

- Replace $\mathbf{u}_i \cdot \mathbf{v}_j$ by $\sum_{g \in n\text{-grams}(w_i)} \mathbf{u}_g \cdot \mathbf{v}_j$

- More to come! Contextualized word embeddings



(Bojanowski et al, 2017): Enriching Word Vectors with Subword Information

Trained word embeddings available

- word2vec: <https://code.google.com/archive/p/word2vec/>
- GloVe: <https://nlp.stanford.edu/projects/glove/>
- FastText: <https://fasttext.cc/>

Download pre-trained word vectors

- Pre-trained word vectors. This data is made available under the [Public Domain Dedication and License](http://www.opendatacommons.org/licenses/pddl/1.0/) v1.0 whose full text can be found at: <http://www.opendatacommons.org/licenses/pddl/1.0/>.
 - [Wikipedia 2014](#) + [Gigaword 5](#) (6B tokens, 400K vocab, uncased, 50d, 100d, 200d, & 300d vectors, 822 MB download): [glove.6B.zip](#)
 - Common Crawl (42B tokens, 1.9M vocab, uncased, 300d vectors, 1.75 GB download): [glove.42B.300d.zip](#)
 - Common Crawl (840B tokens, 2.2M vocab, cased, 300d vectors, 2.03 GB download): [glove.840B.300d.zip](#)
 - Twitter (2B tweets, 27B tokens, 1.2M vocab, uncased, 25d, 50d, 100d, & 200d vectors, 1.42 GB download): [glove.twitter.27B.zip](#)
- Ruby [script](#) for preprocessing Twitter data


Differ in algorithms, text corpora, dimensions, cased/uncased...

Evaluating Word Embeddings

Extrinsic vs intrinsic evaluation

Extrinsic evaluation

- Let's plug these word embeddings into a real NLP system and see whether this improves performance
- Could take a long time but still the most important evaluation metric



Intrinsic evaluation

- Evaluate on a specific/intermediate subtask
- Fast to compute
- Not clear if it really helps the downstream task

Intrinsic evaluation

Word similarity

Example dataset: wordsim-353

353 pairs of words with human judgement

<http://www.cs.technion.ac.il/~gabr/resources/data/wordsim353/>

Word 1	Word 2	Human (mean)
tiger	cat	7.35
tiger	tiger	10
book	paper	7.46
computer	internet	7.58
plane	car	5.77
professor	doctor	6.62
stock	phone	1.62
stock	CD	1.31
stock	jaguar	0.92

Cosine similarity:

$$\cos(\mathbf{u}_i, \mathbf{u}_j) = \frac{\mathbf{u}_i \cdot \mathbf{u}_j}{\|\mathbf{u}_i\|_2 \times \|\mathbf{u}_j\|_2}.$$

Metric: Spearman rank correlation
(Pearson correlation of ranks)

$$r_s = \rho_{R(X), R(Y)} = \frac{\text{cov}(R(X), R(Y))}{\sigma_{R(X)} \sigma_{R(Y)}},$$

Intrinsic evaluation

Word Similarity

Model	Size	WS353	MC	RG	SCWS	RW
SVD	6B	35.3	35.1	42.5	38.3	25.6
SVD-S	6B	56.5	71.5	71.0	53.6	34.7
SVD-L	6B	65.7	<u>72.7</u>	75.1	56.5	37.0
CBOW [†]	6B	57.2	65.6	68.2	57.0	32.5
SG [†]	6B	62.8	65.2	69.7	<u>58.1</u>	37.2
GloVe	6B	<u>65.8</u>	<u>72.7</u>	<u>77.8</u>	53.9	<u>38.1</u>
SVD-L	42B	74.0	76.4	74.1	58.3	39.9
GloVe	42B	<u>75.9</u>	<u>83.6</u>	<u>82.9</u>	<u>59.6</u>	<u>47.8</u>
CBOW*	100B	68.4	79.6	75.4	59.4	45.5

- Spearman rank correlation of word vector similarities with different human judgements on different datasets
- All vectors are 300-dimensional

(Pennington et al, 2014): GloVe: Global Vectors for Word Representation



Intrinsic evaluation

Word analogy

Analogy: Embeddings capture relational meaning!

$\text{vector}('king') - \text{vector}('man') + \text{vector}('woman') \approx \text{vector}('queen')$

$\text{vector}('Paris') - \text{vector}('France') + \text{vector}('Italy') \approx \text{vector}('Rome')$



Intrinsic evaluation

Word analogy

man: woman $\approx$ king: ?

$$\arg \max_i (\cos(\mathbf{u}_i, \mathbf{u}_b - \mathbf{u}_a + \mathbf{u}_c))$$

semantic

Chicago: Illinois $\approx$ Philadelphia: ?

syntactic

bad: worst $\approx$ cool: ?

More examples at

<http://download.tensorflow.org/data/questions-words.txt>

Model	Dim.	Size	Sem.	Syn.	Tot.
ivLBL	100	1.5B	55.9	50.1	53.2
HPCA	100	1.6B	4.2	16.4	10.8
GloVe	100	1.6B	<u>67.5</u>	<u>54.3</u>	<u>60.3</u>
SG	300	1B	61	61	61
CBOW	300	1.6B	16.1	52.6	36.1
vLBL	300	1.5B	54.2	<u>64.8</u>	60.0
ivLBL	300	1.5B	65.2	63.0	64.0
GloVe	300	1.6B	<u>80.8</u>	61.5	<u>70.3</u>
SVD	300	6B	6.3	8.1	7.3
SVD-S	300	6B	36.7	46.6	42.1
SVD-L	300	6B	56.6	63.0	60.1
CBOW [†]	300	6B	63.6	<u>67.4</u>	65.7
SG [†]	300	6B	73.0	66.0	69.1
GloVe	300	6B	<u>77.4</u>	67.0	<u>71.7</u>
CBOW	1000	6B	57.3	68.9	63.7
SG	1000	6B	66.1	65.1	65.6
SVD-L	300	42B	38.4	58.2	49.2
GloVe	300	42B	<u>81.9</u>	<u>69.3</u>	<u>75.0</u>

Hyperparameters and settings matter

Word Similarity								Analogies	
win	Method	WordSim Similarity	WordSim Relatedness	Bruni et al. MEN	Radinsky et al. M. Turk	Luong et al. Rare Words	Hill et al. SimLex	Google Add / Mul	MSR Add / Mul
2	PPMI	.732	.699	.744	.654	.457	.382	.552 / .677	.306 / .535
	SVD	.772	.671	.777	.647	.508	.425	.554 / .591	.408 / .468
	SGNS	.789	.675	.773	.661	.449	.433	.676 / .689	.617 / .644
	GloVe	.720	.605	.728	.606	.389	.388	.649 / .666	.540 / .591
5	PPMI	.732	.706	.738	.668	.442	.360	.518 / .649	.277 / .467
	SVD	.764	.679	.776	.639	.499	.416	.532 / .569	.369 / .424
	SGNS	.772	.690	.772	.663	.454	.403	.692 / .714	.605 / .645
	GloVe	.745	.617	.746	.631	.416	.389	.700 / .712	.541 / .599
10	PPMI	.735	.701	.741	.663	.235	.336	.532 / .605	.249 / .353
	SVD	.766	.681	.770	.628	.312	.419	.526 / .562	.356 / .406
	SGNS	.794	.700	.775	.678	.281	.422	.694 / .710	.520 / .557
	GloVe	.746	.643	.754	.616	.266	.375	.702 / .712	.463 / .519
10	SGNS-LS	.766	.681	.781	.689	.451	.414	.739 / .758	.690 / .729
	GloVe-LS	.678	.624	.752	.639	.361	.371	.732 / .750	.628 / .685

[Improving Distributional Similarity with Lessons Learned from Word Embeddings, Levy, Goldberg, and Dagan, 2015]

Beyond Simple Word Embeddings

Extensions to word vectors

- Subword embeddings (FastText)
- Phrases and multi-word expressions
- Sense embeddings
- Embeddings using other types of context, spaces
- Cross-lingual and cross-modal embeddings
- Context dependent embeddings (Elmo, BERT)

Simple sentence embeddings from word embeddings

- Take average

$$v_s = \frac{1}{|s|} \sum_{w \in s} v_w$$

- Take weighted average

$$v_s = \frac{1}{|s|} \sum_{w \in s} \frac{a}{a + p(w)} v_w$$


Weight rare words more

- Shift sentence embedding

Let V_s be matrix whose columns are the sentence vectors v_s ,
and u be the first singular vector of V_s

$$v'_s = v_s - uu^\top v_s$$

Basic idea: try to subtract out vector corresponding to syntax
(assume it is vector associated with first singular value)




All these are BOW models

Later we will look at models (RNNs)
that take order into account.

[A simple but tough-to-beat baseline for sentence embeddings, Arora, Liang and Ma, 2017]

Sense Embeddings




Improving Word Representations via Global Context and Multiple Word Prototypes [Huang et al, 2012]



Beyond Euclidean spaces


Tree with equally spaced nodes



- Hyperbolic (Poincare) embeddings
- Distances grow exponentially as points go toward the boundary
 - better for hierarchically organized

Poincaré Embeddings for Learning Hierarchical Representations
[Nickel and Kiela, 2017]

Syntax as context





WORD	CONTEXTS
australian	scientist/amod ⁻¹
scientist	australian/amod, discovers/nsubj ⁻¹
discovers	scientist/nsubj, star/dobj, telescope/prep_with
star	discovers/dobj ⁻¹
telescope	discovers/prep_with ⁻¹

Use dependency parse neighbors as context

Dependency based word embeddings
[Levy and Goldberg, 2014]

Cross-lingual word embeddings



See <https://ruder.io/cross-lingual-embeddings/>

Bilingual Word Representations with Monolingual Quality in Mind
[Luong et al, 2015]

Cross-modal embeddings




Figure 5: PCA projection of the 300-dimensional word and image representations for (a) cars and colors and (b) weather and temperature.

Summary

- Word embeddings can be learned using Word2Vec and GloVe
- You can also use SVD on PPMI co-occurrence matrix to obtain dense word vectors
- There is connection between Word2Vec and PPMI co-occurrence matrix
- Evaluating embeddings
- Bias in embeddings
- Some ways to extend embeddings