



Natural Language Processing

Angel Xuan Chang

angelxuanchang.github.io/nlp-class

adapted from lecture slides from

Anoop Sarkar, Danqi Chen and Karthik Narasimhan

Simon Fraser University

Naïve Bayes and Logistic Regression

Naïve Bayes

- Generative Model

$$\hat{c} = \operatorname{argmax}_c P(c)(d|c)$$

- Features assumed to be independent

Logistic Regression

- Discriminative Model

$$\hat{c} = \operatorname{argmax}_c P(c|d)$$

- Features don't have to be independent

Logistic Regression Summary

- Input **features**: $f(x) \rightarrow [f_1, f_2, \dots, f_m]$
- Output: estimate $P(y = c|x)$ for each class c
 - Need to model $P(y = c|x)$ with a family of functions
- Train phase: Learn parameters of model to minimize loss function
 - Need **Loss function** and **Optimization algorithm**
- Test phase: Apply parameters to predict class given a new input

Binary Logistic Regression

- Input features: $f(x) \rightarrow [f_1, f_2, \dots, f_m]$
- Output: $P(y = 1|x)$ and $P(y = 0|x)$
- Classification function: $\sigma(z) = \frac{1}{1+e^{-z}}$
 $z = \mathbf{v} \cdot \mathbf{f}(\mathbf{x})$


bias term

Example ↓

Features: $[1, \text{count}(\text{"amazing"}), \text{count}(\text{"horrible"}), \dots]$

Weights: $[-1.0, 0.8, -0.4, \dots]$

Sigmoid



Learning the weights

- Goal: predict label \hat{y} as close as possible to actual label y
- Distance metric/Loss function: $L(\hat{y}, y)$
- Maximum likelihood estimate:
Choose parameters so that $\log P(y|x)$ is maximized over the training dataset

$$\text{Maximize } \log \prod_{i=1}^n P(y^{(i)} | x^{(i)})$$

where $(x^{(i)}, y^{(i)})$ are paired documents and labels

Binary Cross Entropy Loss

- Let $\hat{y} = \sigma(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}))$
- Classifier probability: $P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$

$$y = 1: P(y|x) = \hat{y} \qquad y = 0: P(y|x) = 1 - \hat{y}$$

- Log probability: $\log P(y|x) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$

Binary Cross Entropy Loss

- Let $\hat{y} = \sigma(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}))$
- Classifier probability: $P(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$
- Log probability: $\log P(y|x) = y \log \hat{y} + (1 - y) \log(1 - \hat{y})$
- Loss:

$$\begin{aligned} L(\hat{y}, y) &= -\log \prod_{i=1}^n P(y^{(i)} | x^{(i)}) = -\sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) \\ &= -\sum_{i=1}^n [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})] \end{aligned}$$

Cross-entropy between the true distribution $P(y|x)$ and predicted distribution $P(\hat{y}|x)$

Binary Cross Entropy Loss

- Cross Entropy Loss:

$$L_{\text{CE}} = - \sum_{i=1}^n \log[y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

- Ranges from 0 (perfect predictions) to $+\infty$
- Lower loss = better classifier

Multinomial Logistic Regression

- Input features: $f(x) \rightarrow [f_1, f_2, \dots, f_m]$
- Output: $P(y = c|x)$ for each class c
- Classification function **Softmax**

$$\frac{\exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}, y))}{\underbrace{\sum_{y'} \exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}, y'))}_{\text{Normalization}}}$$

Features are a function of both input x and output class c

Multinomial Logistic Regression

Var	Definition	Wt
$f_1(0, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	-4.5
$f_1(+, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	2.6
$f_1(-, x)$	$\begin{cases} 1 & \text{if “!”} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1.3

Multinomial Logistic Regression

- Generalize binary loss to multinomial CE loss

$$\begin{aligned} L_{\text{CE}}(\hat{y}, y) &= - \sum_{c=1}^k 1\{y = c\} \log P(y = c | x) \\ &= \sum_{c=1}^k 1\{y = c\} \frac{\exp(\mathbf{v}_c \cdot \mathbf{f}(\mathbf{x}, c))}{\sum_{y'=1}^k \exp(\mathbf{v}_{y'} \cdot \mathbf{f}(\mathbf{x}, y'))} \end{aligned}$$

Multinomial Logistic Regression

- Generalize binary loss to multinomial CE loss

$$\begin{aligned} L_{\text{CE}}(\hat{y}, y) &= - \sum_{c=1}^k 1\{y = c\} \log P(y = c | x) \\ &= \sum_{c=1}^k 1\{y = c\} \frac{\exp(\mathbf{v}_c \cdot \mathbf{f}(\mathbf{x}, c))}{\sum_{y'=1}^k \exp(\mathbf{v}_{y'} \cdot \mathbf{f}(\mathbf{x}, y'))} \end{aligned}$$


Optimization

- We have our loss function and our estimator $\hat{y} = \sigma(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}))$
- How do we find the best set of parameters/weights: \mathbf{v}

$$\hat{\mathbf{v}} = \hat{\theta} = \arg \min \frac{1}{n} \sum_{i=1}^n L_{CE}(y^{(i)}, x^{(i)}; \theta)$$

- Use gradient descent!
 - Find direction of steepest slope
 - Move in opposite direction


Gradient descent (1-D)



$$\theta^{t+1} = \theta^t - \eta \frac{d}{d\theta} f(x; \theta)$$


Gradient descent for LR

- Cross entropy loss for logistic regression is **convex** (i.e. has only one global minimum)
 - No local minima to get stuck in
- Deep neural networks are not so easy
 - Non-convex




Learning Rate

- Updates: $\theta^{t+1} = \theta^t - \eta \frac{d}{d\theta} f(x; \theta)$ Magnitude of movement
- Higher/faster learning rate = larger update



Too small: converge
very slowly




Too big: overshoot and
even diverge

Gradient descent with vector weights

Express slope as a partial derivative of loss w.r.t each weight:

$$\nabla_{\theta} L(f(x; \theta), y) = \begin{bmatrix} \frac{\partial}{\partial w_1} L(f(x; \theta), y) \\ \frac{\partial}{\partial w_2} L(f(x; \theta), y) \\ \vdots \\ \frac{\partial}{\partial w_n} L(f(x; \theta), y) \end{bmatrix}$$



Updates: $\theta^{(t+1)} = \theta^t - \eta \nabla L(f(x; \theta), y)$

Computing the gradients


- From last lecture:

$$\arg \max \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}; \theta)$$

$$\begin{aligned} & \left. \frac{dL(\mathbf{v})}{d\mathbf{v}} \right|_{\mathbf{v}} \\ &= \sum_i \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \frac{1}{\sum_{y''} \exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}_i, y''))} \\ & \quad \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \cdot \exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}_i, y')) \\ &= \sum_i \mathbf{f}(\mathbf{x}_i, y_i) - \sum_i \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \frac{\exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}_i, y'))}{\sum_{y''} \exp(\mathbf{v} \cdot \mathbf{f}(\mathbf{x}_i, y''))} \\ &= \underbrace{\sum_i \mathbf{f}(\mathbf{x}_i, y_i)}_{\text{Observed counts}} - \underbrace{\sum_i \sum_{y'} \mathbf{f}(\mathbf{x}_i, y') \Pr(y' | \mathbf{x}_i; \mathbf{v})}_{\text{Expected counts}} \end{aligned}$$

Stochastic Gradient Descent

- Online optimization
- Compute loss and minimize after each training examples (or mini-batch)



Regularization

- May overfit on the training data!
- Use regularization to prevent overfitting!
- Objective function:

$$\hat{\theta} = \arg \max \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha R(\theta)$$

L2 Regularization

$$R(\theta) = ||\theta||^2 = \sum_{j=1}^d \theta_j^2$$

Euclidean distance of weight vector θ from origin

L2 regularized objective:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha \sum_{j=1}^d \theta_j^2$$

L1 Regularization

$$R(\theta) = ||\theta||_1 = \sum_{j=1}^d |\theta_j|$$

Manhattan distance of weight vector θ from origin

L1 regularized objective:

$$\hat{\theta} = \arg \max_{\theta} \sum_{i=1}^n \log P(y^{(i)} | x^{(i)}) - \alpha \sum_{j=1}^d |\theta_j|$$

L2 vs L1 regularization

- L2 is easier to optimize
 - L1 is complex since the derivative of $|\theta|$ is not continuous at 0
- L2 leads to many small weights
 - L1 prefers sparse weight vectors with many weights set to 0

